

AFR-DT-Copy1

September 19, 2018

```
In [1]: #main libraries
import sqlite3
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings("ignore")

In [2]: #vectorizers
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.feature_extraction.text import CountVectorizer
import gensim
from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

In [3]: #store values in pickles
from sklearn.externals import joblib

In [4]: #performance metrics
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import accuracy_score
from sklearn.metrics import f1_score
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score

In [5]: #modules for building ML model
from sklearn import preprocessing
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import RandomizedSearchCV
from sklearn.model_selection import TimeSeriesSplit
```

0.1 Objective

1. Train, CV, Test split.
2. find right max depth of tree using gridsearchcv().
3. Build Decision tree with featurisation techniques AVGW2V, TFIDFW2V
4. get accuracy, precision scores, confusion matrix, recall score, f1 score.

0.2 Constraints

1. Decision Tree doesn't work well on BOW, TFIDF, also it takes too much time because of dimensionality problem
2. these are solved in AVGW2V, TFIDFW2V where dimensionality is small.
3. only 100k points are used.

```
In [6]: #connect sql database
        con = sqlite3.connect('final.sqlite')

In [7]: #read sql data using pandas
        data = pd.read_sql("SELECT * FROM REVIEWS", con)

In [8]: def partition(x) :
        if x == 'positive' :
            return 1
        return 0

        actualscore = data['Score']
        positivenegative = actualscore.map(partition)
        data['Score'] = positivenegative
```

```
In [9]: data.head()
```

```
Out[9]:
```

	index	Id	ProductId	UserId	ProfileName	\
0	138706	150524	0006641040	ACITT7DI6IDDL	shari zychinski	
1	138688	150506	0006641040	A2IW4PEEK02R0U	Tracy	
2	138689	150507	0006641040	A1S4A3IQ2MU7V4	sally sue "sally sue"	
3	138690	150508	0006641040	AZGXZ2UUK6X	Catherine Hallberg "(Kate)"	
4	138691	150509	0006641040	A3CMRKGE0P909G	Teresa	

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	\
0	0	0	1	939340800	
1	1	1	1	1194739200	
2	1	1	1	1191456000	
3	1	1	1	1076025600	
4	3	4	1	1018396800	

Summary \

```

0          EVERY book is educational
1 Love the book, miss the hard cover version
2          chicken soup with rice months
3      a good swingy rhythm for reading aloud
4          A great way to learn the months

```

Text \

```

0 this witty little book makes my son laugh at l...
1 I grew up reading these Sendak books, and watc...
2 This is a fun way for children to learn their ...
3 This is a great little book to read aloud- it ...
4 This is a book of poetry about the months of t...

```

CleanedText

```

0 witti littl book make son laugh loud recit car...
1 grew read sendak book watch realli rosi movi i...
2 fun way children learn month year learn poem t...
3 great littl book read nice rhythm well good re...
4 book poetri month year goe month cute littl po...

```

```

In [10]: print ('Number of positive & negative data points are \n',data['Score'].value_counts(
        data['Score'].value_counts().plot(kind='bar')

```


Number of positive & negative data points are

1 307061

0 57110

Name: Score, dtype: int64

Out[10]: <matplotlib.axes._subplots.AxesSubplot at 0x1f797e74748>



```
In [11]: #sort data based on time
df_time_sorted = data.sort_values('Time', kind='quicksort')
```

```
In [12]: df_time_sorted.head()
```

```
Out[12]:
```

	index	Id	ProductId	UserId	ProfileName \
	0	138706	150524	0006641040	ACITT7DI6IDDL shari zychinski
	30	138683	150501	0006641040	AJ46FKXOVC7NR Nicholas A Mesiano
	424	417839	451856	B00004CXX9	AIUWLEQ1ADEC5 Elizabeth Medina
	330	346055	374359	B00004CI84	A344SMIA5JECGM Vincent P. Ross
	423	417838	451855	B00004CXX9	AJH6LUC1UT10N The Phantom of the Opera

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time \
0	0	0	1	939340800
30	2	2	1	940809600
424	0	0	1	944092800
330	1	2	1	944438400
423	0	0	1	946857600

	Summary \
0	EVERY book is educational
30	This whole series is great way to spend time w...
424	Entertainingl Funny!
330	A modern day fairy tale
423	FANTASTIC!

	Text \
0	this witty little book makes my son laugh at l...
30	I can remember seeing the show when it aired o...
424	Beetlejuice is a well written movie ... ever...
330	A twist of rumplestiskin captured on film, sta...
423	Beetlejuice is an excellent and funny movie. K...

	CleanedText
0	witti littl book make son laugh loud recit car...
30	rememb see show air televis year ago child sis...
424	beetlejuic well written movi everyth excel act...
330	twist rumplestiskin captur film star michael k...
423	beetlejuic excel funni movi keaton hilari wack...

The important piece of information from dataset for building ML models are text reviews and their Scores if they are positive or negative so lets seperate only those two columns into a seperate dataframe using pandas

```
In [13]: df = pd.DataFrame(data, columns=['CleanedText', 'Score'])
df.head()
```

```
Out[13]:
```

	CleanedText	Score
0	witti littl book make son laugh loud recit car...	1
1	grew read sendak book watch realli rosi movi i...	1
2	fun way children learn month year learn poem t...	1
3	great littl book read nice rhythm well good re...	1
4	book poetri month year goe month cute littl po...	1


```
In [14]: #lets check the total dataset values
df.shape
```

```
Out[14]: (364171, 2)
```

```
In [15]: df_sample = df.head(100000)
print ('Number of +ve & -ve datapoints \n' ,df_sample['Score'].value_counts())
df_sample['Score'].value_counts().plot(kind='bar')
```

```
Number of +ve & -ve datapoints
1    85197
0    14803
Name: Score, dtype: int64
```

```
Out[15]: <matplotlib.axes._subplots.AxesSubplot at 0x1f79df32908>
```



```
In [16]: X = df_sample['CleanedText']
y = df_sample['Score']
print (X.shape)
print (y.shape)
```

```
(100000,)  
(100000,)
```

```
In [17]: #test-train-split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, shuffle=False)  
print('X_train shape :', X_train.shape)  
print('y_train shape :', y_train.shape)  
print('X_test shape :', X_test.shape)  
print('y_test shape :', y_test.shape)
```

```
X_train shape : (70000,)  
y_train shape : (70000,)  
X_test shape : (30000,)  
y_test shape : (30000,)
```

```
In [18]: joblib.dump(X_train, 'X_train.pkl')  
         joblib.dump(X_test, 'X_test.pkl')  
         joblib.dump(X_train, 'y_train.pkl')  
         joblib.dump(X_test, 'y_test.pkl')
```

```
Out[18]: ['y_test.pkl']
```

```
In [102]: X_train = joblib.load('X_train.pkl')  
          X_test = joblib.load('X_test.pkl')  
          y_train = joblib.load('y_train.pkl')  
          y_test = joblib.load('y_test.pkl')
```

0.2.1 check if rows are not shuffled since its time series data

```
In [18]: X_train.head()
```

```
Out[18]: 0    witti littl book make son laugh loud recit car...  
        1    grew read sendak book watch realli rosi movi i...  
        2    fun way children learn month year learn poem t...  
        3    great littl book read nice rhythm well good re...  
        4    book poetri month year goe month cute littl po...  
        Name: CleanedText, dtype: object
```

```
In [19]: X_test.head()
```

```
Out[19]: 70000    introduc madhava agav sister back jan diabet r...  
        70001    love nectar wish amazon would quit rais price ...  
        70002    purchas particular item twice price local heal...  
        70003    madhava agav nectar low kalori natur kosher sw...  
        70004    bought replac honey tendenc crystal winter mon...  
        Name: CleanedText, dtype: object
```

```
In [20]: X_train.tail()
```

```
Out [20]: 69995    madhava agav nectar amber bottl pack use agav ...
69996    forget aspartam artifici sweetner agav nectar ...
69997    ferment agav nectar realli refresh drink twist...
69998    love stuff liquid dissolv easier low gci proba...
69999    start eat healthier one ago biggest step chang...
Name: CleanedText, dtype: object
```

```
In [21]: X_test.tail()
```

```
Out [21]: 99995    delici sugar pretti light brown color delici a...
99996    sugar raw flavor profil much better white suga...
99997    use buy sugar year eat much sugar still sugar ...
99998    product exact advertis save least half retail ...
99999    love sugar also get muscavado sugar great use ...
Name: CleanedText, dtype: object
```

1 Functions to find Hyperparameter & Use Logistic Regression

```
In [22]: def DT_best_params (X_train, y_train) :
# c=1/lambda, lambda = 0.001,0.002,0.01,0.02,0.1,0.2,1,2,10,20,100,200,1000,2000,10000
# gamma = 1/sigma, sigma = 0.001,0.002,0.01,0.02,0.1,0.2,1,2,10,20,100,200,1000,2000,10000
clf = DecisionTreeClassifier()
param_grid = {'max_depth' : list(range(1,50))}
cv = TimeSeriesSplit(n_splits=10)
grid_cv = GridSearchCV(clf, param_grid, cv=cv, verbose=1, n_jobs=-1)
grid_cv.fit(X_train, y_train)
print('best Accuracy:', grid_cv.best_params_)
print('best Score:', grid_cv.best_score_)
# plot depth vs CV error
plt.figure(figsize=(10,6))
plt.plot(param_grid['max_depth'], grid_cv.cv_results_['mean_test_score'], color='g')
plt.xlabel('MAX-DEPTH')
plt.ylabel('Accuracy')
```

```
In [23]: def DT(max_depth, X_train, y_train, X_test, y_test) :
clf = DecisionTreeClassifier(max_depth = max_depth)
clf.fit(X_train,y_train)
y_pred = clf.predict(X_test)
print('accuracy_score =', accuracy_score(y_test, y_pred))
print('precision_score =', precision_score(y_test, y_pred))
print('recall_score =', recall_score(y_test, y_pred))
print('F1_score =', f1_score(y_test, y_pred))
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt="d")
return y_pred
```

2 WORD2VECTOR Model

AVGW2V & TFIDFW2V takes lot of time to train so we use only first 25k data

```
In [24]: # Train your own Word2Vec model using your own text corpus
```

```
i=0
list_of_sent=[]
for sent in X_train.values:
    list_of_sent.append(sent.split())
```

```
In [25]: print(X_train.values[0])
```

```
print("*****")
print(list_of_sent[0])
```

witti littl book make son laugh loud recit car drive along alway sing refrain hes learn whale

['witti', 'littl', 'book', 'make', 'son', 'laugh', 'loud', 'recit', 'car', 'drive', 'along', 'al

```
In [26]: # min_count = 5 considers only words that occurred atleast 5 times
```

```
w2v_model=Word2Vec(list_of_sent,min_count=5,size=50, workers=4)
```

```
In [29]: joblib.dump(w2v_model, 'w2v.pkl')
```

```
Out[29]: ['w2v.pkl']
```

```
In [30]: w2v_model = joblib.load('w2v.pkl')
```

```
In [27]: w2v_words = list(w2v_model.wv.vocab)
```

```
print("number of words that occurred minimum 5 times ",len(w2v_words))
```

```
print("sample words ", w2v_words[0:50])
```

number of words that occurred minimum 5 times 10848

sample words ['littl', 'book', 'make', 'son', 'laugh', 'loud', 'recit', 'car', 'drive', 'along

3 AVGW2V

3.0.1 AVGW2V on train data

```
In [28]: # average Word2Vec
```

```
# compute average word2vec for each review.
```

```
%time train_vectors = []; # the avg-w2v for each sentence/review is stored in this li
```

```
for sent in list_of_sent: # for each review/sentence
```

```
    sent_vec = np.zeros(50) # as word vectors are of zero length
```

```
    cnt_words = 0; # num of words with a valid vector in the sentence/review
```

```
    for word in sent: # for each word in a review/sentence
```

```
        if word in w2v_words:
```

```
            vec = w2v_model.wv[word]
```



```

        sent_vec += vec
        cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    train_vectors.append(sent_vec)
print(len(train_vectors))
print(len(train_vectors[0]))

```

Wall time: 0 ns
70000
50

In [29]: avgw2v_train = preprocessing.normalize(train_vectors)

3.0.2 AVGW2V on test data

In [30]: *# Train your own Word2Vec model using your own test corpus*

```

i=0
list_of_sent_in_test=[]
for sent in X_test.values:
    list_of_sent_in_test.append(sent.split())

```

In [31]: print(X_test.values[0])
print("*****")
print(list_of_sent_in_test[0])

introduc madhava agav sister back jan diabet run famili decid use tea coffe cereal cold hot pa

['introduc', 'madhava', 'agav', 'sister', 'back', 'jan', 'diabet', 'run', 'famili', 'decid', 'u

In [32]: *# average Word2Vec*

compute average word2vec for each review.
test_vectors = []; *# the avg-w2v for each sentence/review is stored in this list*

```

for sent in list_of_sent_in_test : # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    test_vectors.append(sent_vec)
print(len(test_vectors))
print(len(test_vectors[0]))

```

30000
50

```
In [33]: avgw2v_test = preprocessing.normalize(test_vectors)
```

```
In [36]: joblib.dump(avgw2v_train, 'avgw2v_train.pkl')  
         joblib.dump(avgw2v_test, 'avgw2v_test.pkl')
```

```
Out[36]: ['avgw2v_test.pkl']
```


```
In [37]: avgw2v_train = joblib.load('avgw2v_train.pkl')  
         avgw2v_test = joblib.load('avgw2v_test.pkl')
```

```
In [34]: DT_best_params(avgw2v_train, y_train)
```

Fitting 10 folds for each of 49 candidates, totalling 490 fits

```
[Parallel(n_jobs=-1)]: Done 26 tasks      | elapsed: 11.7s  
[Parallel(n_jobs=-1)]: Done 176 tasks     | elapsed: 1.2min  
[Parallel(n_jobs=-1)]: Done 426 tasks     | elapsed: 4.2min  
[Parallel(n_jobs=-1)]: Done 490 out of 490 | elapsed: 5.1min finished
```

best Accuracy: {'max_depth': 6}
best Score: 0.86000314317146




```
In [78]: from sklearn import tree
         tree.export_graphviz(clf,out_file='tree.dot')
```

```
In [35]: DT(6, avgw2v_train, y_train, avgw2v_test, y_test)
```

```
accuracy_score = 0.8648
precision_score = 0.8820290575823239
recall_score = 0.971249755525132
F1_score = 0.9244917715392061
```

```
Out[35]: array([1, 1, 1, ..., 0, 1, 1], dtype=int64)
```



4 TFIDFW2V

4.0.1 TFIDFW2V on Train data

```
In [38]: #calculate TFIDF
```

```
tf_idf_vect = TfidfVectorizer()
final_tf_idf_train = tf_idf_vect.fit_transform(X_train.values)
final_tf_idf_test = tf_idf_vect.transform(X_test.values)
```

```
In [39]: # TF-IDF weighted Word2Vec
```

```
tfidf_feat = tf_idf_vect.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf
np.seterr(divide='ignore', invalid='ignore')
```

```

tfidf_train_vectors = []; # the tfidf-w2v for each sentence/review is stored in this
row=0;
for sent in list_of_sent: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        try:
            vec = w2v_model.wv[word]
            # obtain the tf_idfidf of a word in a sentence/review
            tf_idf = final_tf_idf_train[row, tfidf_feat.index(word)]
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
        except:
            pass
    sent_vec /= weight_sum
    tfidf_train_vectors.append(sent_vec)
    row += 1
print(len(tfidf_train_vectors))
print(len(tfidf_train_vectors[0]))

```

70000

50

```

In [40]: tfidf_w2v_train = preprocessing.normalize(tfidf_train_vectors)
         #tfidf_w2v_train = tfidf_train_vectors

```

4.0.2 TFIDFW2V on Test Data

```

In [41]: # TF-IDF weighted Word2Vec
tfidf_feat = tf_idf_vect.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf
np.seterr(divide='ignore', invalid='ignore')
tfidf_test_vectors = []; # the tfidf-w2v for each sentence/review is stored in this l
row=0;
for sent in list_of_sent_in_test: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        try:
            vec = w2v_model.wv[word]
            # obtain the tf_idfidf of a word in a sentence/review
            tf_idf = final_tf_idf_test[row, tfidf_feat.index(word)]
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
        except:
            pass
    sent_vec /= weight_sum

```

```

        tfidf_test_vectors.append(sent_vec)
        row += 1
    print(len(tfidf_test_vectors))
    print(len(tfidf_test_vectors[0]))

30000
50

In [42]: tfidf2v_test = preprocessing.normalize(tfidf_test_vectors)
        tfidf2v_test.shape

Out[42]: (30000, 50)

In [45]: joblib.dump(tfidf2v_train, 'tfidf2v_train.pkl')
        joblib.dump(tfidf2v_test, 'tfidf2v_test.pkl')

Out[45]: ['tfidf2v_test.pkl']

In [46]: tfidf2v_train = joblib.load('tfidf2v_train.pkl')
        tfidf2v_test = joblib.load('tfidf2v_test.pkl')

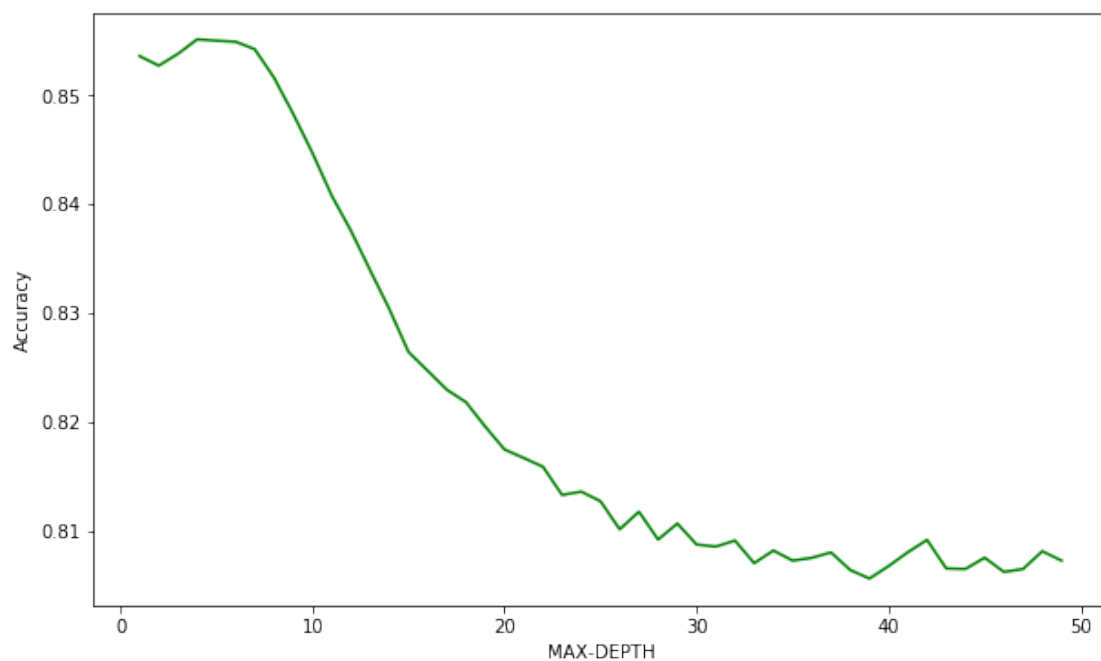
In [43]: DT_best_params(tfidf2v_train, y_train)

Fitting 10 folds for each of 49 candidates, totalling 490 fits

[Parallel(n_jobs=-1)]: Done 26 tasks      | elapsed: 11.6s
[Parallel(n_jobs=-1)]: Done 176 tasks     | elapsed: 1.2min
[Parallel(n_jobs=-1)]: Done 426 tasks     | elapsed: 4.0min
[Parallel(n_jobs=-1)]: Done 490 out of 490 | elapsed: 4.9min finished

best Accuracy: {'max_depth': 4}
best Score: 0.8551783749803552


```



```
In [44]: DT(4, tfidf2v_train, y_train, tfidf2v_test, y_test)
```

```
accuracy_score = 0.8551  
precision_score = 0.8715676660128888  
recall_score = 0.9734011343633875  
F1_score = 0.919674039580908
```

```
Out[44]: array([1, 1, 1, ..., 1, 1, 1], dtype=int64)
```



5 BAG of WORDS

```
In [47]: vect = CountVectorizer()
```

```
In [48]: bow_X_train = vect.fit_transform(X_train)
bow_X_train = preprocessing.normalize(bow_X_train)
bow_X_train
```

```
Out[48]: <70000x32149 sparse matrix of type '<class 'numpy.float64'>'
with 2162199 stored elements in Compressed Sparse Row format>
```

```
In [49]: bow_X_test = vect.transform(X_test)
bow_X_test = preprocessing.normalize(bow_X_test)
bow_X_test
```

```
Out[49]: <30000x32149 sparse matrix of type '<class 'numpy.float64'>'
with 880827 stored elements in Compressed Sparse Row format>
```

```
In [ ]: joblib.dump(bow_X_train, 'bow_X_train.pkl')
joblib.dump(bow_X_test, 'bow_X_test.pkl')
```

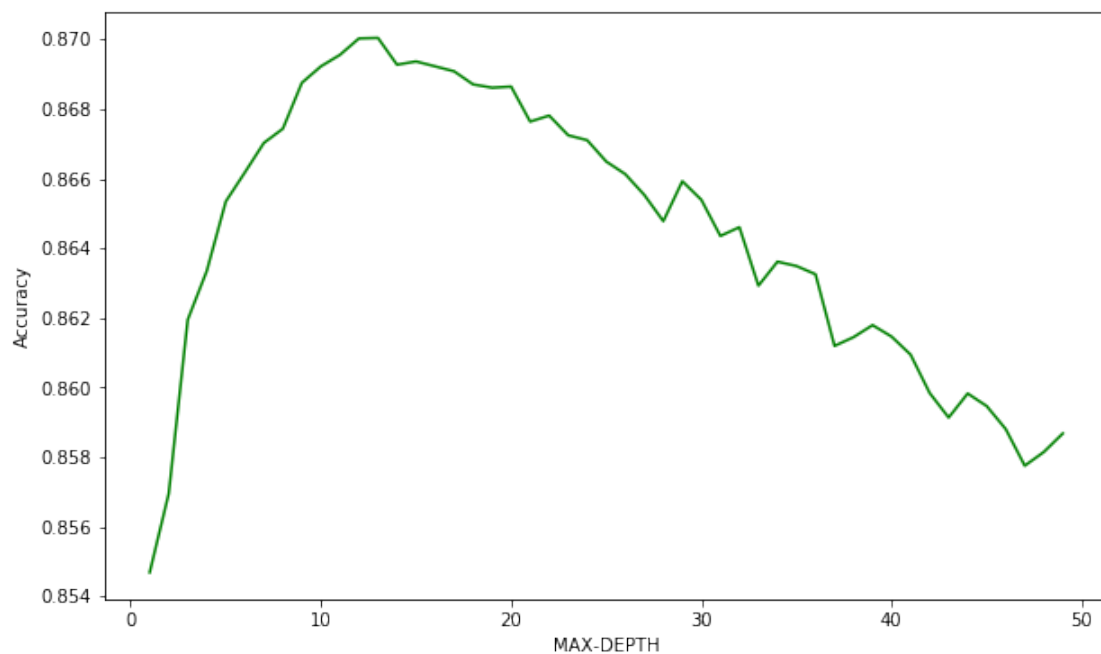
```
In [ ]: bow_X_train = joblib.load('bow_X_train.pkl')
bow_X_test = joblib.load('bow_X_test.pkl')
```

```
In [50]: DT_best_params(bow_X_train, y_train)
```

Fitting 10 folds for each of 49 candidates, totalling 490 fits

```
[Parallel(n_jobs=-1)]: Done 26 tasks      | elapsed: 12.5s  
[Parallel(n_jobs=-1)]: Done 176 tasks     | elapsed: 1.4min  
[Parallel(n_jobs=-1)]: Done 426 tasks     | elapsed: 7.5min  
[Parallel(n_jobs=-1)]: Done 490 out of 490 | elapsed: 10.0min finished
```


best Accuracy: {'max_depth': 13}
best Score: 0.87004557598617



```
In [51]: %time DT(13, bow_X_train, y_train, bow_X_test, y_test)
```

```
accuracy_score = 0.8722  
precision_score = 0.8813392763134805  
recall_score = 0.9822804615685508  
F1_score = 0.9290761774390469  
Wall time: 6.55 s
```

```
Out[51]: array([1, 1, 1, ..., 1, 1, 1], dtype=int64)
```

6 TFIDF

```
In [20]: vect = TfidfVectorizer()
```

```
In [21]: from sklearn import preprocessing
tfidf_X_train = vect.fit_transform(X_train)
tfidf_X_train = preprocessing.normalize(tfidf_X_train)
tfidf_X_train
```

```
Out[21]: <70000x32149 sparse matrix of type '<class 'numpy.float64'>'
         with 2162199 stored elements in Compressed Sparse Row format>
```

```
In [22]: tfidf_X_test = vect.transform(X_test)
tfidf_X_test = preprocessing.normalize(tfidf_X_test)
tfidf_X_test
```

```
Out[22]: <30000x32149 sparse matrix of type '<class 'numpy.float64'>'
         with 880827 stored elements in Compressed Sparse Row format>
```

```
In [109]: joblib.dump(tfidf_X_train, 'tfidf_X_train.pkl')
          joblib.dump(tfidf_X_test, 'tfidf_X_test.pkl')
```

```
Out[109]: ['tfidf_X_test.pkl']
```

```
In [111]: tfidf_X_train = joblib.load('tfidf_X_train.pkl')
          tfidf_X_test = joblib.load('tfidf_X_test.pkl')
```


```
In [56]: DT_best_params(tfidf_X_train, y_train)
```

Fitting 10 folds for each of 49 candidates, totalling 490 fits

```
[Parallel(n_jobs=-1)]: Done 26 tasks      | elapsed: 7.8s  
[Parallel(n_jobs=-1)]: Done 176 tasks     | elapsed: 1.1min  
[Parallel(n_jobs=-1)]: Done 426 tasks     | elapsed: 6.0min  
[Parallel(n_jobs=-1)]: Done 490 out of 490 | elapsed: 7.9min finished
```

best Accuracy: {'max_depth': 13}

best Score: 0.8687254439729687



```
In [57]: DT(13, tfidf_X_train, y_train, tfidf_X_test, y_test)
```


```
accuracy_score = 0.8706666666666667
```

```
precision_score = 0.882141472526698
```

```
recall_score = 0.9790338353217289
```

```
F1_score = 0.928065556750343
```

```
Out[57]: array([1, 1, 1, ..., 1, 1, 1], dtype=int64)
```




7 Lets see the feature importance and decision tree on our best model TFIDF

```
In [18]: def DT(max_depth, X_train, y_train, X_test, y_test) :
          clf = DecisionTreeClassifier(max_depth = max_depth)
          clf.fit(X_train,y_train)
          y_pred = clf.predict(X_test)
          print('accuracy_score =', accuracy_score(y_test, y_pred))
          print('precision_score =', precision_score(y_test, y_pred))
          print('recall_score =', recall_score(y_test, y_pred))
          print('F1_score =', f1_score(y_test, y_pred))
          cm = confusion_matrix(y_test, y_pred)
          sns.heatmap(cm, annot=True, fmt="d")
          fi = clf.feature_importances_
          return y_pred, fi
```

```
In [23]: tfidf-fi = DT(13, tfidf_X_train, y_train, tfidf_X_test, y_test)
```

```
accuracy_score = 0.8704666666666667
precision_score = 0.8821165438713998
recall_score = 0.9787991394484648
F1_score = 0.9279463027516132
```



```
In [24]: features = vect.get_feature_names()
```

```
In [29]: tfidf-fi[1]
```

```
Out[29]: array([0., 0., 0., ..., 0., 0., 0.])
```


```
In [35]: # Plotting word cloud
from wordcloud import WordCloud

freq = tfidf-fi[1]
words = vect.get_feature_names()
#freqs = tag_dtm.sum(axis=0).A1
result = dict(zip(words, freq))

# Lets first convert the 'result' dictionary to 'list of tuples'
tup = dict(result.items())
#Initializing WordCloud using frequencies of tags.
wordcloud = WordCloud(
    background_color='black',
    width=1600,
    height=800,
).generate_from_frequencies(tup)

fig = plt.figure(figsize=(30,20))
plt.imshow(wordcloud)
plt.axis('off')
plt.tight_layout(pad=0)
```

```
fig.savefig("tag.png")
plt.show()
```



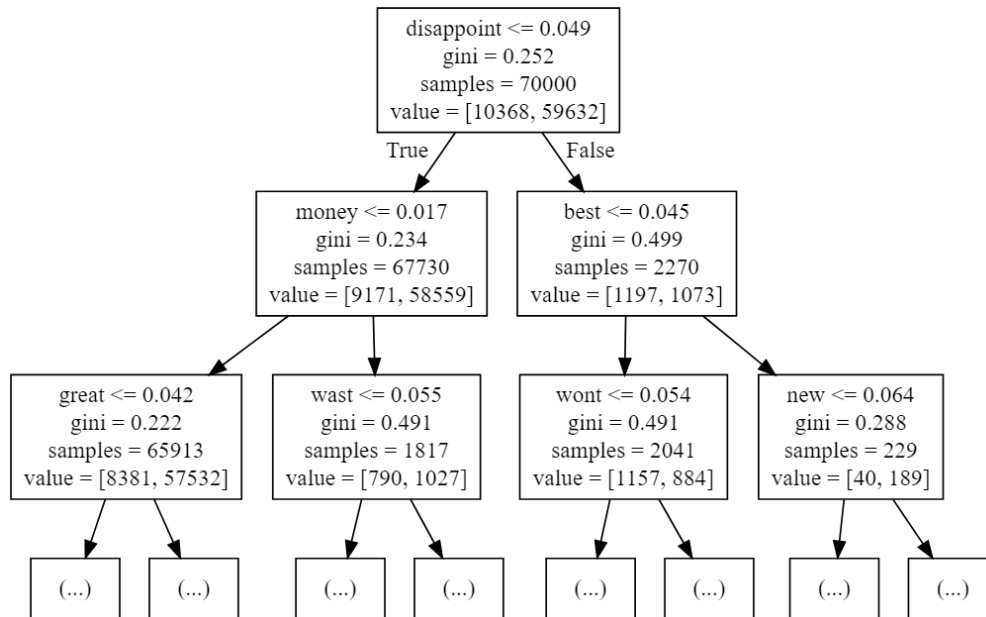
```
In [88]: clf = DecisionTreeClassifier(max_depth = 13)
         clf.fit(tfidf_X_train, y_train)
```

```
Out[88]: DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=13,
                                max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, presort=False, random_state=None,
                                splitter='best')
```

```
In [107]: from sklearn import tree
          tree = tree.export_graphviz(clf, out_file='TFIDF_D tree.dot', max_depth=2, feature_names=feature_names)
```

8 RESULTS

```
In [58]: from prettytable import PrettyTable
x = PrettyTable()
x.field_names = ["MODEL", "Tree max-depth", "ACCURACY", "PRECISION", "RECALL", 'F1-SCORE']
#BOW
x.add_row(['BOW with DT GridSearch', '13', 0.87, 0.88, 0.98, 0.92])
x.add_row(['--'*5, '--'*5, '--'*8, '--'*5, '--'*5, '--'*5])
#TFIDF
x.add_row(['TFIDF with DT GridSearch', '13', 0.87, 0.88, 0.97, 0.92])
x.add_row(['--'*5, '--'*8, '--'*8, '--'*5, '--'*5, '--'*5])
```



alt text

```

#AVGW2V
x.add_row(['AVGW2V with DT GridSearch', '6', 0.86, 0.88, 0.97, 0.92])
x.add_row(['--'*5, '-'*8, '-'*8, '-'*5, '--'*5, '--'*5])
#TFIDFW2V
x.add_row(['TFIDFW2V with DT GridSearch', '4', 0.85, 0.87, 0.97, 0.91])
print(x)

```

MODEL	Tree max-depth	ACCURACY	PRECISION	RECALL	F1-SCORE
BOW with DT GridSearch	13	0.87	0.88	0.98	0.92
TFIDF with DT GridSearch	13	0.87	0.88	0.97	0.92
AVGW2V with DT GridSearch	6	0.86	0.88	0.97	0.92
TFIDFW2V with DT GridSearch	4	0.85	0.87	0.97	0.91

```

In [108]: #number of positive and negative values in test data
          y_test.value_counts()

```

```

Out[108]: 1    25565
          0     4435
          Name: Score, dtype: int64

```

OBSERVATIONS

since AVGW2v and TFIDFW2V took too much time for converting to a vector. the total number of datapoints used are limited to 100K. also, the BOW & TFIDF were trained on all data and the confusion matrix and accuracy score were same in percentages.

1. The best results were obtained from TFIDF with closely 3.8k mis-classifications out of 30k datapoints.
2. DT still lags behind Logistic Regression Model. the maximum misclassification in logistic regression was 2.5k. but this exceeds to atleast 4k misclassification in all models.
3. almost all the model till now has tough in predicting negative reviews.
4. since this data is imbalanced, there are large amount of data for positive reviews. so, the False positive rate is very high on almost all the vectorizers. W2V perform very bad on this.
5. using SKlearn DT .dot file is generated and sample DT is visualised online graphviz and as we can see in DT. Disappoint is the main word with highest gini impurity