

# AFR-SVC

September 19, 2018

```
In [1]: #main libraries
import sqlite3
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings("ignore")

In [2]: #vectorizers
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.feature_extraction.text import CountVectorizer
import gensim
from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

In [3]: #store values in pickles
from sklearn.externals import joblib

In [4]: #performance metrics
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import accuracy_score
from sklearn.metrics import f1_score
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score

In [5]: #modules for building ML model
from sklearn import preprocessing
from sklearn.model_selection import train_test_split
from sklearn.linear_model import SGDClassifier
from sklearn.svm import SVC
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import RandomizedSearchCV
from sklearn.model_selection import TimeSeriesSplit
```

## 0.1 Objective

1. Train, CV, Test split.
2. find right 'c' (1/Lambda), gamma(1/sigma) using gridsearchcv(), randomsearchcv().
3. Build SVC with featurisation techniques like BOW, TFIDF AVGW2V2 TFIDFW2V and use l1 or l2 regularizer.
4. get accuracy, precision scores, confusion matrix, recall score, f1 score.

```
In [8]: #connect sql database
con = sqlite3.connect('final.sqlite')
```

```
In [9]: #read sql data using pandas
data = pd.read_sql("SELECT * FROM REVIEWS", con)
```

```
In [10]: def partition(x) :
          if x == 'positive' :
              return 1
          return 0

          actualscore = data['Score']
          positivenegative = actualscore.map(partition)
          data['Score'] = positivenegative
```

```
In [11]: data.head()
```

```
Out[11]:
```

	index	Id	ProductId	UserId	ProfileName
0	138706	150524	0006641040	ACITT7DI6IDDL	shari zychinski
1	138688	150506	0006641040	A2IW4PEEK02ROU	Tracy
2	138689	150507	0006641040	A1S4A3IQ2MU7V4	sally sue "sally sue"
3	138690	150508	0006641040	AZGXZ2UUK6X	Catherine Hallberg "(Kate)"
4	138691	150509	0006641040	A3CMRKGEOP909G	Teresa

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time
0	0	0	1	939340800
1	1	1	1	1194739200
2	1	1	1	1191456000
3	1	1	1	1076025600
4	3	4	1	1018396800

	Summary
0	EVERY book is educational
1	Love the book, miss the hard cover version
2	chicken soup with rice months
3	a good swingy rhythm for reading aloud
4	A great way to learn the months

Text \

```

0  this witty little book makes my son laugh at l...
1  I grew up reading these Sendak books, and watc...
2  This is a fun way for children to learn their ...
3  This is a great little book to read aloud- it ...
4  This is a book of poetry about the months of t...

```

CleanedText

```

0  witty littl book make son laugh loud recit car...
1  grew read sendak book watch realli rosi movi i...
2  fun way children learn month year learn poem t...
3  great littl book read nice rhythm well good re...
4  book poetri month year goe month cute littl po...

```

```

In [12]: print ('Number of positive & negative data points are \n',data['Score'].value_counts(
        data['Score'].value_counts().plot(kind='bar')

```


Number of positive & negative data points are

1 307061

0 57110

Name: Score, dtype: int64

Out[12]: <matplotlib.axes.\_subplots.AxesSubplot at 0x27f30075588>



```

In [13]: #sort data based on time
df_time_sorted = data.sort_values('Time', kind='quicksort')

```

```
In [14]: df_time_sorted.head()
```

```
Out[14]:
```

	index	Id	ProductId	UserId	ProfileName	\
0	138706	150524	0006641040	ACITT7DI6IDDL	shari zychinski	
30	138683	150501	0006641040	AJ46FKXOVC7NR	Nicholas A Mesiano	
424	417839	451856	B00004CXX9	AIUWLEQ1ADEG5	Elizabeth Medina	
330	346055	374359	B00004CI84	A344SMIA5JECGM	Vincent P. Ross	
423	417838	451855	B00004CXX9	AJH6LUC1UT1ON	The Phantom of the Opera	

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	\
0	0	0	1	939340800	
30	2	2	1	940809600	
424	0	0	1	944092800	
330	1	2	1	944438400	
423	0	0	1	946857600	

	Summary	\
0	EVERY book is educational	
30	This whole series is great way to spend time w...	
424	Entertainingl Funny!	
330	A modern day fairy tale	
423	FANTASTIC!	

	Text	\
0	this witty little book makes my son laugh at l...	
30	I can remember seeing the show when it aired o...	
424	Beetlejuice is a well written movie ... ever...	
330	A twist of rumplestiskin captured on film, sta...	
423	Beetlejuice is an excellent and funny movie. K...	

	CleanedText
0	witti littl book make son laugh loud recit car...
30	rememb see show air televis year ago child sis...
424	beetlejuic well written movi everyth excel act...
330	twist rumplestiskin captur film star michael k...
423	beetlejuic excel funni movi keaton hilari wack...

The important piece of information from dataset for building ML models are text reviews and their Scores if they are positive or negative so lets seperate only those two columns into a seperate dataframe using pandas

```
In [15]: df = pd.DataFrame(data, columns=['CleanedText', 'Score'])
df.head()
```

```
Out[15]:
```

	CleanedText	Score
0	witti littl book make son laugh loud recit car...	1
1	grew read sendak book watch realli rosi movi i...	1
2	fun way children learn month year learn poem t...	1

```

3 great littl book read nice rhythm well good re... 1
4 book poetri month year goe month cute littl po... 1

```

```

In [16]: #lets check the total dataset values
df.shape

```

```

Out[16]: (364171, 2)

```

```

In [15]: df_sample = df
print ('Number of +ve & -ve datapoints \n',df_sample['Score'].value_counts())
df_sample['Score'].value_counts().plot(kind='bar')

```

```

Number of +ve & -ve datapoints

```

```

1    307061

```

```

0     57110

```

```


Name: Score, dtype: int64

```

```

Out[15]: <matplotlib.axes._subplots.AxesSubplot at 0x2665f1c3e10>

```



```

In [16]: X = df_sample['CleanedText']
y = df_sample['Score']
print (X.shape)
print (y.shape)

```

```

(364171,)

```

```

(364171,)

```

```
In [17]: #test-train-split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, shuffle=False)
print('X_train shape :', X_train.shape)
print('y_train shape :', y_train.shape)
print('X_test shape :', X_test.shape)
print('y_test shape :', y_test.shape)
```

```
X_train shape : (254919,)
y_train shape : (254919,)
X_test shape : (109252,)
y_test shape : (109252,)
```

```
In [18]: joblib.dump(X_train, 'X_train.pkl')
joblib.dump(X_test, 'X_test.pkl')
joblib.dump(X_train, 'y_train.pkl')
joblib.dump(X_test, 'y_test.pkl')
```

```
Out[18]: ['y_test.pkl']
```

```
In [19]: X_train = joblib.load('X_train.pkl')
X_test = joblib.load('X_test.pkl')
y_train = joblib.load('y_train.pkl')
y_test = joblib.load('y_test.pkl')
```

### 0.1.1 check if rows are not shuffled since its time series data

```
In [18]: X_train.head()
```

```
Out[18]: 0    witti littl book make son laugh loud recit car...
1    grew read sendak book watch realli rosi movi i...
2    fun way children learn month year learn poem t...
3    great littl book read nice rhythm well good re...
4    book poetri month year goe month cute littl po...
Name: CleanedText, dtype: object
```

```
In [19]: X_test.head()
```

```
Out[19]: 254919    word want reduc caffien brand best choic use ...
254920    louisiana nativ like peopl louisiana drink com...
254921    drink number brand coffe definit dont enjoy fi...
254922    purchas great price amazon pleas flavor buddi ...
254923    love coffe still block form fresh communiti co...
Name: CleanedText, dtype: object
```

```
In [20]: X_train.tail()
```

```
Out[20]: 254914    havent found decaff serv challah one give rebb
254915    purchas coffe base posit feedback either got b...
```

```

254916    drank communiti coffe mani year recent becam c...
254917    bought pack give other kept one tri glad great...
254918    love communiti coffe yummi strong without grea...
Name: CleanedText, dtype: object

```

```
In [21]: X_test.tail()
```

```

Out[21]: 364166    love love sweeten use bake unsweeten flavor co...
364167    tri sauc believ start littl sweet honey tast b...
364168    bought hazelnut past nocciola spread local sho...
364169    purchas product local store kid love quick eas...
364170    purchas send son whos away colleg deliv right ...
Name: CleanedText, dtype: object

```

Since SVM is Computationally expensive and very time consuming. and specifically GridSearch takes a lot of time for SVC(). so, we use - SGDClassifier with hinge loss along with GridSearchCV & RandomizedSearchCV

## 1 Functions to find Hyperparameter & Use Logistic Regression

```

In [22]: def SGD_best_params (X_train, y_train) :
          #c=1/lambda, lambda = 0.001,0.002,0.01,0.02,0.1,0.2,1,2,10,20,100,200,1000,2000,10000
          #gamma = 1/sigma, sigma = 0.001,0.002,0.01,0.02,0.1,0.2,1,2,10,20,100,200,1000,2000,10000
          clf = SGDClassifier(loss='hinge')
          cv= TimeSeriesSplit(n_splits=10)
          param_grid = {'alpha':[1000,500,100,50,10,5,1,0.5,0.1,0.05,0.01,0.005,0.001,0.0001],
                        'penalty':['l1', 'l2', 'elasticnet']}
          grid_cv = GridSearchCV(clf, param_grid, cv=cv, verbose=1, n_jobs=-1)
          grid_cv.fit(X_train,y_train)
          print("Best HyperParameter: ",grid_cv.best_params_)
          print("Best Accuracy: ", (grid_cv.best_score_*100))
          #accessing cv_results
          cv_results = pd.DataFrame(grid_cv.cv_results_)
          alpha_cv = cv_results[['param_alpha', 'mean_test_score']]
          plot_data = alpha_cv.loc[:,3, :]
          #Function for cv_error vs alpha plot
          plt.figure(figsize=(10,6))
          plt.xlabel('Alpha')
          plt.ylabel('Best Score')
          plt.plot(plot_data['param_alpha'], plot_data['mean_test_score'], marker='o', mark

In [23]: def SGD(alpha, penalty, X_train, y_train, X_test, y_test) :
          clf = SGDClassifier(alpha=alpha, penalty=penalty, loss='hinge')
          clf.fit(X_train, y_train)
          y_pred=clf.predict(X_test)
          print('accuracy_score =', accuracy_score(y_test, y_pred))
          print('precision_score =', precision_score(y_test, y_pred))
          print('recall_score =', recall_score(y_test, y_pred))

```

```

        cm = confusion_matrix(y_test, y_pred)
        sns.heatmap(cm, annot=True, fmt="d")
        return y_pred

In [49]: #c=1/lambda, lambda = 0.001,0.002,0.01,0.02,0.1,0.2,1,2,10,20,100,200,1000,2000,10000
def SGD_best_params_rand (X_train, y_train) :
    clf = SGDClassifier(loss='hinge')
    cv= TimeSeriesSplit(n_splits=10)
    #param_grid = {'C':[1000,500,100,50,10,5,1,0.5,0.1,0.05,0.01,0.005,0.001,0.0005,0.0001]}
    param_grid = {'alpha':[1000,500,100,50,10,5,1,0.5,0.1,0.05,0.01,0.005,0.001,0.0001],
                  'penalty':['l1','l2', 'elasticnet']}
    rand_cv = RandomizedSearchCV(clf, param_grid, cv=cv, verbose=1, n_jobs=-1, random_state=42)
    rand_cv.fit(X_train,y_train)
    print("Best HyperParameter: ",rand_cv.best_params_)
    print("Best Accuracy: ", (rand_cv.best_score_*100))
    #accessing cv_results
    cv_results = pd.DataFrame(rand_cv.cv_results_)
    alpha_cv = cv_results[['param_alpha', 'mean_test_score']]
    plot_data = alpha_cv.sort_values('param_alpha')
    #Function for cv_error vs alpha plot
    plt.figure(figsize=(10,6))
    plt.xlabel('Alpha')
    plt.ylabel('Best Score')
    plt.plot(plot_data['param_alpha'], plot_data['mean_test_score'], marker='o', mark

```

```

In [25]: def SGD_rand(alpha, penalty, X_train, y_train, X_test, y_test) :
    #clf = SVC(C=C , gamma=gamma )
    clf = SGDClassifier(alpha=alpha, penalty=penalty, loss='hinge')
    clf.fit(X_train, y_train)
    y_pred = clf.predict(X_test)
    print('accuracy_score =', accuracy_score(y_test, y_pred))
    print('precision_score =', precision_score(y_test, y_pred))
    print('recall_score =', recall_score(y_test, y_pred))
    cm = confusion_matrix(y_test, y_pred)
    sns.heatmap(cm, annot=True, fmt="d")
    return y_pred

```

## 2 BAG of WORDS

```
In [54]: vect = CountVectorizer()
```

```
In [55]: bow_X_train = vect.fit_transform(X_train)
        bow_X_train = preprocessing.normalize(bow_X_train)
        bow_X_train

```

```
Out[55]: <254919x59601 sparse matrix of type '<class 'numpy.float64'>'
        with 7863068 stored elements in Compressed Sparse Row format>
```



```

In [56]: bow_X_test = vect.transform(X_test)
         bow_X_test = preprocessing.normalize(bow_X_test)
         bow_X_test

Out[56]: <109252x59601 sparse matrix of type '<class 'numpy.float64'>'
         with 3581565 stored elements in Compressed Sparse Row format>

In [57]: joblib.dump(bow_X_train, 'bow_X_train.pkl')
         joblib.dump(bow_X_test, 'bow_X_test.pkl')

Out[57]: ['bow_X_test.pkl']

In [58]: bow_X_train = joblib.load('bow_X_train.pkl')
         bow_X_test = joblib.load('bow_X_test.pkl')

In [29]: SGD_best_params(bow_X_train, y_train)

```

Fitting 10 folds for each of 45 candidates, totalling 450 fits


```

[Parallel(n_jobs=-1)]: Done 26 tasks      | elapsed:    8.9s
[Parallel(n_jobs=-1)]: Done 176 tasks     | elapsed:   35.4s
[Parallel(n_jobs=-1)]: Done 426 tasks     | elapsed:   1.4min
[Parallel(n_jobs=-1)]: Done 450 out of 450 | elapsed:   1.5min finished

```

Best HyperParameter: {'alpha': 0.0001, 'penalty': 'l2'}


Best Accuracy: 90.57305601104686



```
In [30]: %time SGD(0.0001, 'l2', bow_X_train, y_train, bow_X_test, y_test)
```

```
accuracy_score = 0.9013839563577783  
precision_score = 0.9016603834122082  
recall_score = 0.9893722357911413  
Wall time: 528 ms
```

```
Out[30]: array([1, 1, 1, ..., 1, 1, 1], dtype=int64)
```




```
In [59]: SGD_best_params_rand(bow_X_train, y_train)
```

Fitting 10 folds for each of 10 candidates, totalling 100 fits

```
[Parallel(n_jobs=-1)]: Done 26 tasks      | elapsed:    9.3s  
[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed:   23.2s finished
```


```
Best HyperParameter: {'penalty': 'l2', 'alpha': 0.0001}  
Best Accuracy: 90.59894709588332
```



```
In [32]: SGD_rand(0.0001, 'l2', bow_X_train, y_train, bow_X_test, y_test)
```

```
accuracy_score = 0.9008256141763995  
precision_score = 0.9007638937556942  
recall_score = 0.9898453143221775
```

```
Out[32]: array([1, 1, 1, ..., 1, 1, 1], dtype=int64)
```



### 3 TFIDF

```
In [33]: vect = TfidfVectorizer()
```

```
In [34]: from sklearn import preprocessing
tfidf_X_train = vect.fit_transform(X_train)
tfidf_X_train = preprocessing.normalize(tfidf_X_train)
tfidf_X_train
```

```
Out[34]: <254919x59601 sparse matrix of type '<class 'numpy.float64'>'
         with 7863068 stored elements in Compressed Sparse Row format>
```

```
In [35]: tfidf_X_test = vect.transform(X_test)
tfidf_X_test = preprocessing.normalize(tfidf_X_test)
tfidf_X_test
```

```
Out[35]: <109252x59601 sparse matrix of type '<class 'numpy.float64'>'
         with 3581565 stored elements in Compressed Sparse Row format>
```

```
In [ ]: joblib.dump(tfidf_X_train, 'tfidf_X_train.pkl')
        joblib.dump(tfidf_X_test, 'tfidf_X_test.pkl')
```

```
In [ ]: tfidf_X_train = joblib.load(tfidf_X_train, 'tfidf_X_train.pkl')
        tfidf_X_test = joblib.load(tfidf_X_test, 'tfidf_X_test.pkl')
```


```
In [36]: SGD_best_params(tfidf_X_train, y_train)
```

Fitting 10 folds for each of 45 candidates, totalling 450 fits

```
[Parallel(n_jobs=-1)]: Done 26 tasks      | elapsed: 8.5s
[Parallel(n_jobs=-1)]: Done 176 tasks     | elapsed: 34.4s
[Parallel(n_jobs=-1)]: Done 426 tasks     | elapsed: 1.4min
[Parallel(n_jobs=-1)]: Done 450 out of 450 | elapsed: 1.4min finished
```

Best HyperParameter: {'alpha': 0.0001, 'penalty': 'l2'}


Best Accuracy: 90.53378786571157



```
In [37]: SGD(0.0001, 'l2', tfidf_X_train, y_train, tfidf_X_test, y_test)
```

```
accuracy_score = 0.8994984073518105
precision_score = 0.8979226004342024
recall_score = 0.9919686667986886
```

```
Out[37]: array([1, 1, 1, ..., 1, 1, 1], dtype=int64)
```




```
In [38]: SGD_best_params_rand(tfidf_X_train, y_train)
```

Fitting 10 folds for each of 10 candidates, totalling 100 fits

```
[Parallel(n_jobs=-1)]: Done 26 tasks      | elapsed:    8.6s  
[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed:   20.5s finished
```

Best HyperParameter: {'penalty': 'l2', 'alpha': 0.0001}

Best Accuracy: 90.59074825235177




```
In [39]: SGD_rand(0.0001, 'l2', tfidf_X_train, y_train, tfidf_X_test, y_test)
```

```
accuracy_score = 0.8993794163951232
```

```
precision_score = 0.897774680141385
```

```
recall_score = 0.9920126741039013
```

```
Out[39]: array([1, 1, 1, ..., 1, 1, 1], dtype=int64)
```



## 4 WORD2VECTOR Model

AVGW2V & TFIDFW2V takes lot of time to train so we use only first 100k data

```
In [26]: w2v_data = df.head(100000)
         print ('Number of +ve & -ve datapoints \n' ,df_sample['Score'].value_counts())
         w2v_data['Score'].value_counts().plot(kind='bar')
```


Number of +ve & -ve datapoints

1 307061

0 57110

Name: Score, dtype: int64

Out[26]: <matplotlib.axes.\_subplots.AxesSubplot at 0x266672328d0>



```
In [27]: w2v_X = w2v_data['CleanedText']
         w2v_y = w2v_data['Score']
         print (w2v_X.shape)
         print (w2v_y.shape)
```

(100000,)

(100000,)



```
In [28]: #test-train-split
w2v_X_train, w2v_X_test, w2v_y_train, w2v_y_test = train_test_split(w2v_X, w2v_y, test_size=0.2, random_state=42)
print('X_train shape : ', w2v_X_train.shape)
print('y_train shape : ', w2v_y_train.shape)
print('X_test shape : ', w2v_X_test.shape)
print('y_test shape : ', w2v_y_test.shape)
```

```
In [29]: # Train your own Word2Vec model using your own text corpus
i=0
list_of_sent=[]
for sent in w2v_X_train.values:
    list_of_sent.append(sent.split())
```

```
witti littl book make son laugh loud recit car drive along alway sing refrain hes learn whale :
*****
['witti', 'littl', 'book', 'make', 'son', 'laugh', 'loud', 'recit', 'car', 'drive', 'along', 'a
```

```
In [ ]: joblib.dump(w2v, 'w2v.pkl')
```

```
number of words that occurred minimum 5 times 10848
sample words ['littl', 'book', 'make', 'son', 'laugh', 'loud', 'recit', 'car', 'drive', 'along']
```

### 5.0.1 AVGW2V on train data

```

%time train_vectors = []; # the avg-w2v for each sentence/review is stored in this li
for sent in list_of_sent: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    train_vectors.append(sent_vec)
print(len(train_vectors))
print(len(train_vectors[0]))

```

Wall time: 0 ns  
254919  
50

In [45]: avgw2v\_train = preprocessing.normalize(train\_vectors)

## 5.0.2 AVGW2V on test data

```

In [39]: # Train your own Word2Vec model using your own text corpus
i=0
list_of_sent_in_test=[]
for sent in w2v_X_test.values:
    list_of_sent_in_test.append(sent.split())

```

```

In [40]: print(w2v_X_test.values[0])
print("*****")
print(list_of_sent_in_test[0])

```

introduc madhava agav sister back jan diabet run famili decid use tea coffe cereal cold hot pa  
\*\*\*\*\*  
['introduc', 'madhava', 'agav', 'sister', 'back', 'jan', 'diabet', 'run', 'famili', 'decid', 'pa']

```

In [48]: # average Word2Vec
# compute average word2vec for each review.
test_vectors = []; # the avg-w2v for each sentence/review is stored in this list

for sent in list_of_sent_in_test : # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]

```

```

        sent_vec += vec
        cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    test_vectors.append(sent_vec)
print(len(test_vectors))
print(len(test_vectors[0]))

```

109252

50

```
In [49]: avgw2v_test = preprocessing.normalize(test_vectors)
```

```
In [ ]: joblib.dump(avgw2v_train, 'avgw2v_train.pkl')
        joblib.dump(avgw2v_test, 'avgw2v_test.pkl')
```

```
In [ ]: avgw2v_train = joblib.load('avgw2v_train.pkl')
        avgw2v_test = joblib.load('avgw2v_test.pkl')
```

```
In [50]: SGD_best_params(avgw2v_train, y_train)
```

Fitting 10 folds for each of 45 candidates, totalling 450 fits


```

[Parallel(n_jobs=-1)]: Done 26 tasks      | elapsed:    9.3s
[Parallel(n_jobs=-1)]: Done 176 tasks     | elapsed:   35.8s
[Parallel(n_jobs=-1)]: Done 426 tasks     | elapsed:   1.4min
[Parallel(n_jobs=-1)]: Done 450 out of 450 | elapsed:   1.5min finished

```

Best HyperParameter: {'alpha': 0.0001, 'penalty': 'l2'}


Best Accuracy: 89.81703633382239



```
In [51]: SGD(0.0001, 'l2', avgw2v_train, y_train, avgw2v_test, y_test)
```

```
accuracy_score = 0.8903086442353458  
precision_score = 0.9058571810644558  
recall_score = 0.968842827909433
```

```
Out [51]: array([1, 1, 1, ..., 1, 1, 1], dtype=int64)
```




```
In [52]: SGD_best_params_rand(avgw2v_train, y_train)
```

Fitting 10 folds for each of 10 candidates, totalling 100 fits

```
[Parallel(n_jobs=-1)]: Done 26 tasks      | elapsed:    8.8s  
[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed:   20.7s finished
```

Best HyperParameter: {'penalty': 'l2', 'alpha': 0.0001}


Best Accuracy: 90.59506343315785



```
In [53]: SGD_rand(0.0001, 'l2', avgw2v_train, y_train, avgw2v_test, y_test)
```

```
accuracy_score = 0.8867938344378135  
precision_score = 0.8938430365525819  
recall_score = 0.9803617400488481
```

```
Out [53]: array([1, 1, 1, ..., 1, 1, 1], dtype=int64)
```



## 6 TFIDFW2V

### 6.0.1 TFIDFW2V on Train data

In [33]: *#calculate TFIDF*

```
tf_idf_vect = TfidfVectorizer()
final_tf_idf_train = tf_idf_vect.fit_transform(w2v_X_train.values)
final_tf_idf_test = tf_idf_vect.transform(w2v_X_test.values)
```

In [34]: *# TF-IDF weighted Word2Vec*

```
tfidf_feat = tf_idf_vect.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf
np.seterr(divide='ignore', invalid='ignore')
tfidf_train_vectors = []; # the tfidf-w2v for each sentence/review is stored in this row=0;
for sent in list_of_sent: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        try:
            vec = w2v_model.wv[word]
            # obtain the tf_idfidf of a word in a sentence/review
            tf_idf = final_tf_idf_train[row, tfidf_feat.index(word)]
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
```

```

        except:
            pass
        sent_vec /= weight_sum
        tfidf_train_vectors.append(sent_vec)
        row += 1
print(len(tfidf_train_vectors))
print(len(tfidf_train_vectors[0]))

```

70000

50

```

In [44]: tfidf2v_train = preprocessing.normalize(tfidf_train_vectors)
         #tfidf2v_train = tfidf_train_vectors

```

## 6.0.2 TFIDFW2V on Test Data

```

In [41]: # TF-IDF weighted Word2Vec
tfidf_feat = tf_idf_vect.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf
np.seterr(divide='ignore', invalid='ignore')
tfidf_test_vectors = []; # the tfidf-w2v for each sentence/review is stored in this l
row=0;
for sent in list_of_sent_in_test: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        try:
            vec = w2v_model.wv[word]
            # obtain the tf_idfidf of a word in a sentence/review
            tf_idf = final_tf_idf_test[row, tfidf_feat.index(word)]
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
        except:
            pass
    sent_vec /= weight_sum
    tfidf_test_vectors.append(sent_vec)
    row += 1
print(len(tfidf_test_vectors))
print(len(tfidf_test_vectors[0]))

```

30000

50

```

In [43]: tfidf2v_test = preprocessing.normalize(tfidf_test_vectors)
         tfidf2v_test.shape

```

Out [43]: (30000, 50)



```
In [ ]: joblib.dump(tfidf2v_train, 'tfidf2v_train.pkl')
        joblib.dump(tfidf2v_test, 'tfidf2v_test.pkl')
```

```
In [ ]: tfidf2v_train = joblib.load('tfidf2v_train.pkl')
        tfidf2v_test = joblib.load('tfidf2v_test.pkl')
```


```
In [46]: SGD_best_params(tfidf2v_train, w2v_y_train)
```

Fitting 10 folds for each of 45 candidates, totalling 450 fits

```
[Parallel(n_jobs=-1)]: Done 26 tasks      | elapsed:    6.8s
[Parallel(n_jobs=-1)]: Done 176 tasks     | elapsed:   16.3s
[Parallel(n_jobs=-1)]: Done 426 tasks     | elapsed:   31.4s
[Parallel(n_jobs=-1)]: Done 450 out of 450 | elapsed:   32.4s finished
```

Best HyperParameter: {'alpha': 0.0001, 'penalty': 'l2'}


Best Accuracy: 87.82806852113782



```
In [47]: SGD(0.0001, 'l2', tfidf2v_train, w2v_y_train, tfidf2v_test, w2v_y_test)
```

```
accuracy_score = 0.8780333333333333
precision_score = 0.883320501154896
recall_score = 0.9872873068648543
```

```
Out[47]: array([1, 1, 1, ..., 1, 1, 1], dtype=int64)
```




```
In [50]: SGD_best_params_rand(tfidf_w2v_train, w2v_y_train)
```

Fitting 10 folds for each of 10 candidates, totalling 100 fits

```
[Parallel(n_jobs=-1)]: Done 26 tasks      | elapsed:    6.6s  
[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed:   11.2s finished
```

Best HyperParameter: {'penalty': 'l2', 'alpha': 0.0001}

Best Accuracy: 87.78249253496779




```
In [51]: SGD_rand(0.0001, 'l2', tfidf2v_train, w2v_y_train, tfidf2v_test, w2v_y_test)
```

```
accuracy_score = 0.8765
```

```
precision_score = 0.8800417246175244
```

```
recall_score = 0.9900254253862703
```

```
Out[51]: array([1, 1, 1, ..., 1, 1, 1], dtype=int64)
```



## 7 SVM WITH SKLEARN CLASSIFIER ON OUR BEST MODEL BOW

```
In [79]: SVM_df_sample = df.head(100000)
         print ('Number of +ve & -ve datapoints \n' ,SVM_df_sample['Score'].value_counts())
         SVM_df_sample['Score'].value_counts().plot(kind='bar')
```


Number of +ve & -ve datapoints

1 85197

0 14803

Name: Score, dtype: int64

Out[79]: <matplotlib.axes.\_subplots.AxesSubplot at 0x27f3ca809e8>



```
In [80]: SVM_X = SVM_df_sample['CleanedText']
         SVM_y = SVM_df_sample['Score']
         print (SVM_X.shape)
         print (SVM_y.shape)
```

(100000,)

(100000,)

```

In [81]: #test-train-split
        SVM_X_train, SVM_X_test, SVM_y_train, SVM_y_test = train_test_split(SVM_X, SVM_y, test_size=0.3, random_state=42)
        print('X_train shape :', SVM_X_train.shape)
        print('y_train shape :', SVM_y_train.shape)
        print('X_test shape :', SVM_X_test.shape)
        print('y_test shape :', SVM_y_test.shape)

X_train shape : (70000,)
y_train shape : (70000,)
X_test shape : (30000,)
y_test shape : (30000,)

In [82]: vect = CountVectorizer()

In [83]: bow_X_train = vect.fit_transform(SVM_X_train)
        bow_X_train = preprocessing.normalize(bow_X_train)
        bow_X_train

Out[83]: <70000x32149 sparse matrix of type '<class 'numpy.float64'>'
        with 2162199 stored elements in Compressed Sparse Row format>

In [84]: bow_X_test = vect.transform(SVM_X_test)
        bow_X_test = preprocessing.normalize(bow_X_test)
        bow_X_test


Out[84]: <30000x32149 sparse matrix of type '<class 'numpy.float64'>'
        with 880827 stored elements in Compressed Sparse Row format>

In [86]: from sklearn.svm import SVC
        clf = SVC(max_iter = 800)
        clf.fit(bow_X_train, SVM_y_train)
        SVM_y_pred=clf.predict(bow_X_test)
        print('accuracy_score =', accuracy_score(SVM_y_test, SVM_y_pred))
        print('precision_score =', precision_score(SVM_y_test, SVM_y_pred))
        print('recall_score =', recall_score(SVM_y_test, SVM_y_pred))
        cm = confusion_matrix(SVM_y_test, SVM_y_pred)
        sns.heatmap(cm, annot=True, fmt="d")

accuracy_score = 0.6995
precision_score = 0.9571318086399293
recall_score = 0.677723450029337

Out[86]: <matplotlib.axes._subplots.AxesSubplot at 0x27f42f5b128>

```



## 8 RESULTS

```
In [88]: from prettytable import PrettyTable
x = PrettyTable()
x.field_names = ["MODEL", "Alpha & PENALTY", "ACCURACY", "PRECISION", "RECALL" ]
#BOW
x.add_row(['BOW with SVM GridSearch', '5 & L2', 0.90, 0.90, 0.99])
x.add_row(["BOW with SVM Random", '1 & L1', 0.90, 0.90, 0.99])
x.add_row(['--'*5, '-'*5, '-'*8, '-'*5, '--'*5])
#TFIDF
x.add_row(['TFIDF with SVM GridSearch', '5 & L2', 0.89, 0.89, 0.99])
x.add_row(["TFIDF with SVM Random", '10 & L1', 0.89, 0.89, 0.99])
x.add_row(['--'*5, '-'*8, '-'*8, '-'*5, '--'*5])
#AVGW2V
x.add_row(['AVGW2V with SVM GridSearch', '100 & L2', 0.89, 0.89, 0.98])
x.add_row(["AVGW2V with SVM Random", '10 & L1', 0.88, 0.88, 0.97])
x.add_row(['--'*5, '-'*8, '-'*8, '-'*5, '--'*5])
#TFIDFW2V
x.add_row(['TFIDFW2V with SVM GridSearch', '0.000 & L2', 0.87, 0.88, 0.98])
x.add_row(["TFIDFW2V with SVM Random", '0.0001 & L2', 0.87, 0.88, 0.99])
x.add_row(['--'*5, '-'*8, '-'*8, '-'*5, '--'*5])
#SVC
x.add_row(['BOW with SVC', 'Default', 0.69, 0.95, 0.67])
print(x)
```

MODEL	Alpha & PENALTY	ACCURACY	PRECISION	RECALL
BOW with SVM GridSearch	5 & L2	0.9	0.9	0.99
BOW with SVM Random	1 & L1	0.9	0.9	0.99
-----	-----	-----	-----	-----
TFIDF with SVM GridSearch	5 & L2	0.89	0.89	0.99
TFIDF with SVM Random	10 & L1	0.89	0.89	0.99
-----	-----	-----	-----	-----
AVGW2V with SVM GridSearch	100 & L2	0.89	0.89	0.98
AVGW2V with SVM Random	10 & L1	0.88	0.88	0.97
-----	-----	-----	-----	-----
TFIDFW2V with SVM GridSearch	0.000 & L2	0.87	0.88	0.98
TFIDFW2V with SVM Random	0.0001 & L2	0.87	0.88	0.99
-----	-----	-----	-----	-----
BOW with SVC	Default	0.69	0.95	0.67

```
In [60]: #number of positive and negative values in test data
         y_test.value_counts()
```

```
Out[60]: 1    90894
         0    18358
         Name: Score, dtype: int64
```

### OBSERVATIONS

since AVGW2v and TFIDFW2V took too much time for converting to a vector. the total number of datapoints used are limited to 100K. also, the BOW & TFIDF were trained on all data and the confusion matrix and accuracy score were same in percentages.

1. The best results were obtained from BOW with closely 10.7k mis-classifications out of 100k datapoints.
2. SVM lags behind Logistic Regression Model. the maximum misclassification in logistic regression was 2.5k. but this exceeds to atleast 3k misclassification in all models.
3. The original SVM classifier in SKlearn is trained on our best model with default values and the best values we get is less than what we achieved in SGDClassifier with hinge loss.
4. SVM gave 10k misclassification out of 100k points than SGD which showed 3k misclassification. so it is good to use SGD with Hingeloss which is fast and also can operate on whole data.
5. the precision and recall were high on almost all the models.
6. Till now the best model on amazon fine food reviews is Logistic Regression it has accuracy of 92 and misclassification of 2.3K reviews.
7. since this data is imbalanced, there are large amount of data for positive reviews. so, the False positive rate is very high on almost all the vectorizers. W2V perform very bad on this.